

# Generative Adversarial Networks

Presented by “GANerators” team

*Marwa Nair, Anna Krysta, Caio Rocha*

01.

# PROBLEM STATEMENT

---

# Problem Statement

## Goal:

Train a GAN on the MNIST dataset, with a fixed generator architecture, and implement techniques to improve the data generation.

## Problems with Vanilla GANs:

- Mode Collapse: The generator produces a limited variety of samples, ignoring parts of the data distribution.
- Vanishing gradients: The discriminator becomes too strong, resulting in poor generator updates.
- Overfitting: The discriminator becomes too focused on the generator's distribution, and fails to generalize.

02.

# TECHNIQUES AND PRELIMINARY RESULTS

---

# During training


Differential Privacy GAN (Alex B., Gautam K. & Guojun Z.; 2023)

- We clip the Discriminator gradients and add Gaussian noise
- Generator is unchanged
- We perform more Discriminator steps per Generator step to preserve balance between the Generator and Discriminator

# During training

## Collaborative Sampling (CS)

- Given a fixed generator  $G$  and a fixed discriminator  $D$ .
- Refines the generated samples through gradient-based updates at the last layer of the generator.
- Shifts the generator distribution closer to the real data distribution. (Liu, Y., Kothari, P., & Alahi, A. ; 2020)



# During training

## Collaborative Sampling (CS)

1. Iteratively refines the generated samples

$$x_l^{k+1} = x_l^k - \lambda \nabla_l \mathcal{L}_G(x_l^k),$$

$$x^{k+1} = G_L \circ G_{L-1} \circ \dots \circ G_l(x_l^{k+1}),$$

2. Minimizes the loss of the generator

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

3. Minimizes the loss of the discriminator

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_r} [\log D(x)] - \mathbb{E}_{x' \sim p_c} [1 - \log D(x')],$$

# During training

Collaborative Sampling (CS)

1. Iteratively refines the generated samples

**Output of layer l**  $\longrightarrow x_l^{k+1} = x_l^k - \lambda \nabla_l \mathcal{L}_G(x_l^k),$

**Output of the last layer**  $\longrightarrow x^{k+1} = G_L \circ G_{L-1} \circ \dots \circ G_l(x_l^{k+1}),$

2. Minimizes the loss of the generator

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

3. Minimizes the loss of the discriminator

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_r} [\log D(x)] - \mathbb{E}_{x' \sim p_c} [1 - \log D(x')],$$



# During training

## Collaborative Sampling (CS)

1. Iteratively refines the generated samples

$$x_l^{k+1} = x_l^k - \lambda \nabla_l \mathcal{L}_G(x_l^k),$$

$$x^{k+1} = G_L \circ G_{L-1} \circ \dots \circ G_l(x_l^{k+1}),$$

2. Minimizes the loss of the generator

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

**Refined sample**






























3. Minimizes the loss of the discriminator

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_r} [\log D(x)] - \mathbb{E}_{x' \sim p_c} [1 - \log D(x')],$$

# Results

## Collaborative Sampling (CS)

	FID	Precision	Recall
CS	26.82	0.53	0.23

Project Name	Image 1	Image 2	Image 3	Image 4	Image 5	Image 6	Image 7	Image 8	Image 9
ganerators									
									
									

# During generation

Latent Space Interpolation (LSI)

- Purpose** Improve GAN output quality by sampling interpolated points in latent space, creating smoother, more realistic transitions.
- Intuition** Blending features in the latent space explores intermediate regions for greater image consistency. (*Radford et al., 2015 ; Karras et al., 2019*)

# During generation

Latent Space Interpolation (LSI)

## Pseudo-algorithm

```
for i in range(nb_images // interpolation_steps):
```

- Sample two random latent vectors  $z_1$  and  $z_2$
- Linearly interpolate between them at different intervals defined by **interpolation\_steps**:

$$z_{\text{interp}} = (1 - \alpha) \cdot z_1 + \alpha \cdot z_2 \quad \text{for } \alpha \in \text{linspace}(0, 1, \text{interpolation\_steps})$$

- Feed each  $z_{\text{interp}}$  into the generator to generate new images.

# During generation

Latent Space Interpolation (LSI)

## Results

	FID	Precision	Recall
VanillaGAN + LSI	21.81	0.57	0.16
CS + LSI	21.35	0.57	0.19

03.

# COMPARISON OF RESULTS

---

# Comparative Analysis

	FID	Precision	Recall
VanillaGAN	26.84	0.53	0.25
CS	26.82	0.53	0.23
VanillaGAN + LSI	21.81	0.57	0.16
CS + LSI	21.35	0.57	0.19

**THANK YOU FOR  
YOUR ATTENTION**

*Any questions?*



# Collaborative sampling algorithms

---

**Algorithm 1** Collaborative Sampling

---

```
1: Input: a frozen generator  $G$ , a frozen discriminator  $D$ ,  
   the layer index for sample refinement  $l$ , the maximum  
   number of steps  $K$ , the stopping criterion  $\eta$   
2: Output: a synthetic sample  $x$   
3: Randomly draw a latent code  $z$   
4:  $x^0 \leftarrow \text{ProposeSample}(G, z)$   
5: for  $k = 0, 1, \dots, K - 1$  do  
6:   if  $D(x^k) < \eta$  then  
7:      $g_l^k \leftarrow \text{GetGradient}(D, x_l^k),$   
8:      $x_l^{k+1} \leftarrow \text{UpdateActivation}(g_l^k, x_l^k), \quad (\text{Eq. 3})$   
9:      $x^{k+1} \leftarrow \text{UpdateSample}(G, x_l^{k+1}), \quad (\text{Eq. 4})$   
10:  else  
11:    break  
12:  end if  
13: end for
```

---

# Collaborative sampling algorithms

---

**Algorithm 2** Discriminator Shaping

---

- 1: **Input:** a frozen generator  $G$ , a pre-trained discriminator  $D$ , the batch size  $m$
  - 2: **Output:** a fine-tuned discriminator  $\tilde{D}$
  - 3: **for** number of D shaping iterations **do**
  - 4:   Draw  $m$  refined samples  $\{x_c^{(1)}, \dots, x_c^{(m)}\}$  from the collaborative data distribution  $p_c(x)$  according to Algorithm 1
  - 5:   Draw  $m$  real samples  $\{x_r^{(1)}, \dots, x_r^{(m)}\}$  from the real data distribution  $p_r(x)$
  - 6:   Shape the discriminator by minimizing the objective function Eq. 6
  - 7: **end for**
-

# Differential privacy GAN

---

**Algorithm 1** TrainDPGAN( $D; \phi_0, \theta_0, \text{OptD}, \text{OptG}, n_{\mathcal{D}}, T, B, C, \sigma, \delta$ )

---

```
1: Input: Labelled dataset  $D = \{(x_j, y_j)\}_{j=1}^n$ . Discriminator  $\mathcal{D}$  and generator  $\mathcal{G}$  initializations  $\phi_0$  and  $\theta_0$ . Opti-
   mizers  $\text{OptD}$ ,  $\text{OptG}$ . Hyperparameters:  $n_{\mathcal{D}}$  ( $\mathcal{D}$  steps per  $\mathcal{G}$  step),  $T$  (total number of  $\mathcal{D}$  steps),  $B$  (expected batch
   size),  $C$  (clipping norm), and  $\sigma$  (noise level). Privacy parameter  $\delta$ .
2:  $q \leftarrow B/|D|$  and  $t, k \leftarrow 0$  ▷ Calculate sampling rate  $q$ , initialize counters.
3: while  $t < T$  do ▷ Update  $\mathcal{D}$  with DPSGD.
4:    $S_t \sim \text{PoissonSample}(D, q)$  ▷ Sample a real batch  $S_t$  by including each  $(x, y) \in D$  w.p.  $q$ .
5:    $\tilde{S}_t \sim \mathcal{G}(\cdot; \theta_k)^B$  ▷ Sample fake batch  $\tilde{S}_t$ .
6:    $g_{\phi_t} \leftarrow \sum_{(x,y) \in S_t} \text{clip}(\nabla_{\phi_t}(-\log(\mathcal{D}(x, y; \phi_t))); C)$ 
      $+ \sum_{(\tilde{x}, \tilde{y}) \in \tilde{S}_t} \text{clip}(\nabla_{\phi_t}(-\log(1 - \mathcal{D}(\tilde{x}, \tilde{y}; \phi_t))); C)$  ▷ Clip per-example gradients.
7:    $\hat{g}_{\phi_t} \leftarrow \frac{1}{2B}(g_{\phi_t} + z_t)$ , where  $z_t \sim \mathcal{N}(0, C^2 \sigma^2 I)$  ▷ Add Gaussian noise.
8:    $\phi_{t+1} \leftarrow \text{OptD}(\phi_t, \hat{g}_{\phi_t})$ 
9:    $t \leftarrow t + 1$ 
10:  if  $n_{\mathcal{D}}$  divides  $t$  then ▷ Perform  $\mathcal{G}$  update every  $n_{\mathcal{D}}$  steps.
11:     $\tilde{S}'_t \sim \mathcal{G}(\cdot; \theta_k)^B$ 
12:     $g_{\theta_k} \leftarrow \frac{1}{B} \sum_{(\tilde{x}, \tilde{y}) \in \tilde{S}'_t} \nabla_{\theta_k}(-\log(\mathcal{D}(\tilde{x}, \tilde{y}; \phi_t)))$ 
13:     $\theta_{k+1} \leftarrow \text{OptG}(\theta_k, g_{\theta_k})$ 
14:     $k \leftarrow k + 1$ 
15:  end if
16: end while
17:  $\varepsilon \leftarrow \text{PrivacyAccountant}(T, \sigma, q, \delta)$  ▷ Compute privacy budget spent.
18: Output: Final  $\mathcal{G}$  parameters  $\theta_k$  and  $(\varepsilon, \delta)$ -DP guarantee.
```

---